

Running Apache Kafka at scale on Kubernetes with Strimzi

May 2024



Speakers



Steffen Karlsson
Lead Software Engineer
Streaming Services, Maersk

steffen.karlsson@maersk.com



Maciej Tatarski
Software Engineer
RETinA Team, Maersk

maciej.tatarski@maersk.com

Agenda

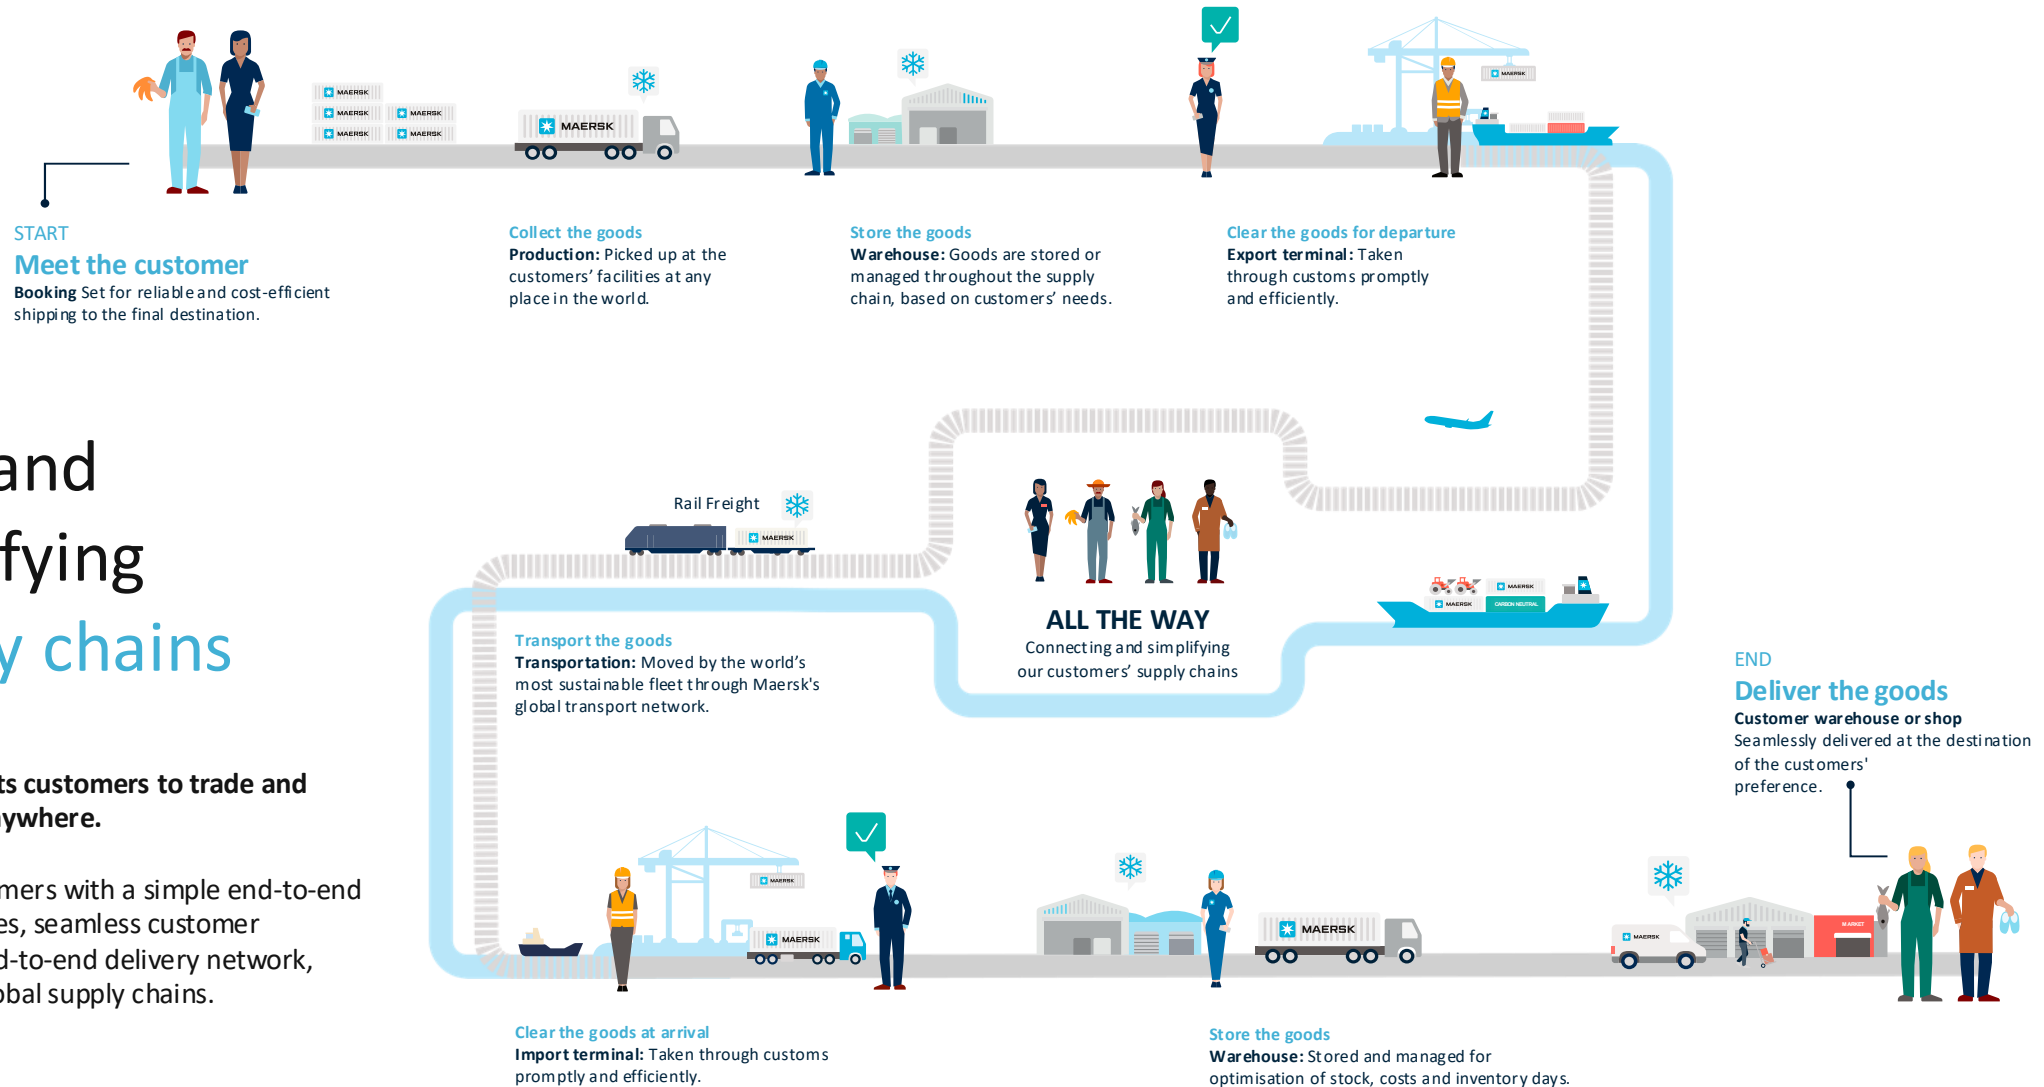
- What is Maersk?
- Backstory
- DevX Vision
- Technology backbone of Maersk
- Running Kafka on Kubernetes?
- Architecture
- On-Prem / Edge
- Tools
- Observability & Monitoring
- Alerting
- Stream Processing
- Reflections



Connecting and simplifying global supply chains

A.P. Moller - Maersk enables its customers to trade and grow by transporting goods anywhere.

Maersk works to provide customers with a simple end-to-end offering of products and services, seamless customer engagement and a superior end-to-end delivery network, taking the complexity out of global supply chains.



Backstory



Strategy to move from batch to event driven enterprise architecture



Fast and easy bootstrap with PaaS and SaaS



Open-source adoption strategy, benefitting from the economy of scale



Building up internal capabilities and inner-sourcing for a more tailored solution



Transition



Gain trust of the business



Cost-benefit analysis:
Pay-as-you-go

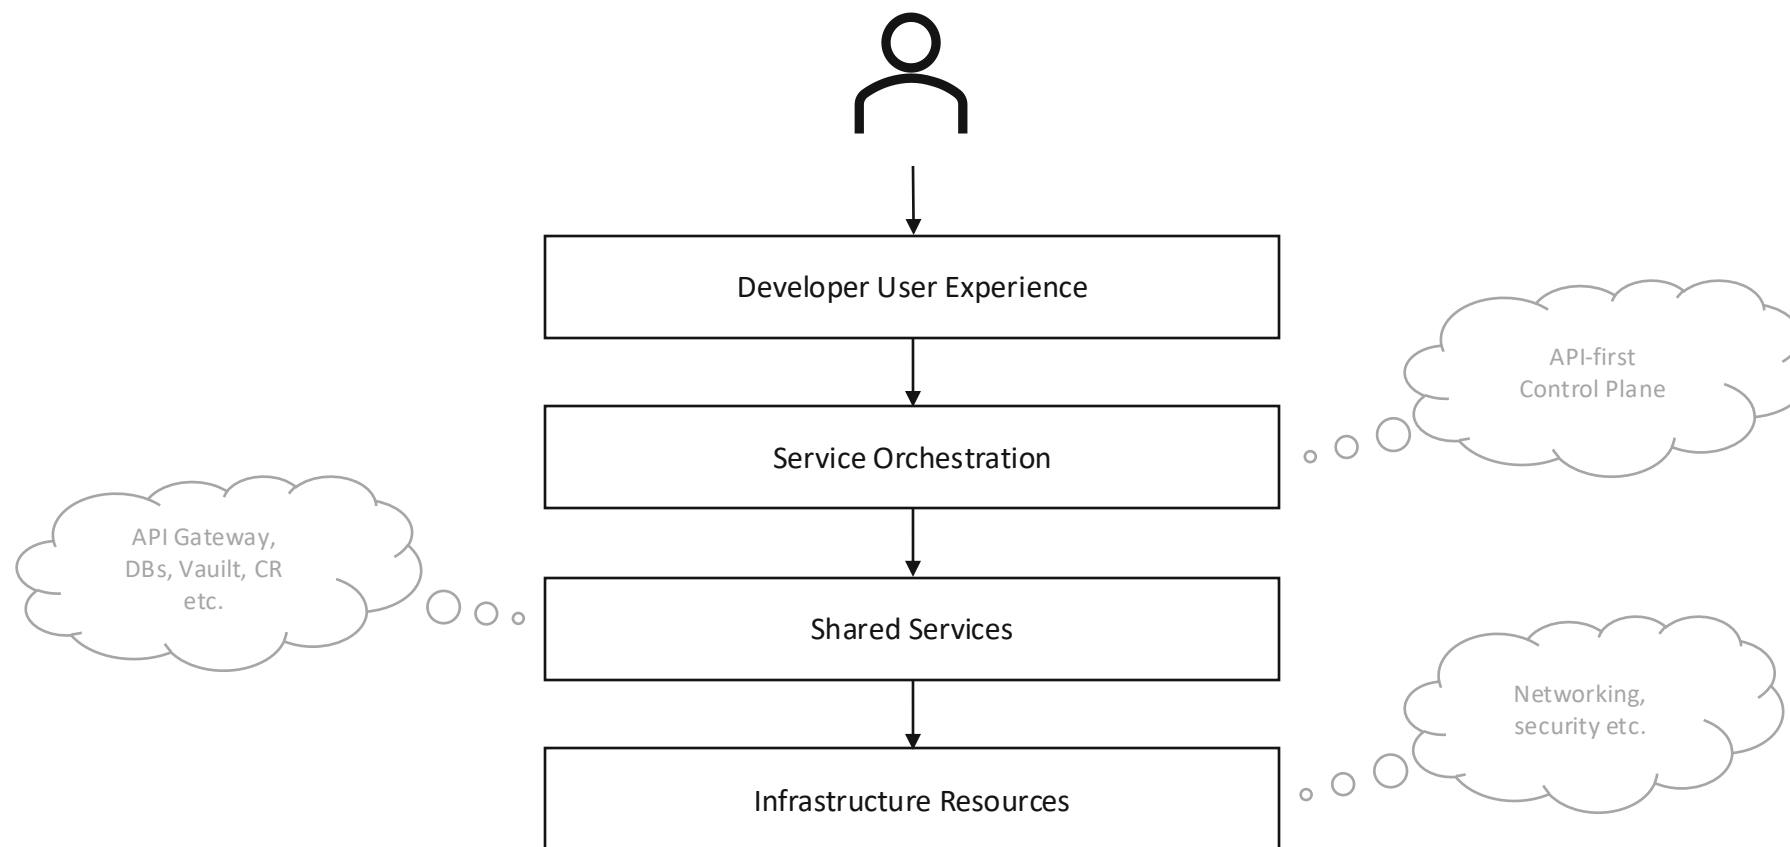


Internal knowledge on how to manage and administer a Kafka ecosystem




Rapidly increasing adoption of the platform, point of no return

DevX Vision

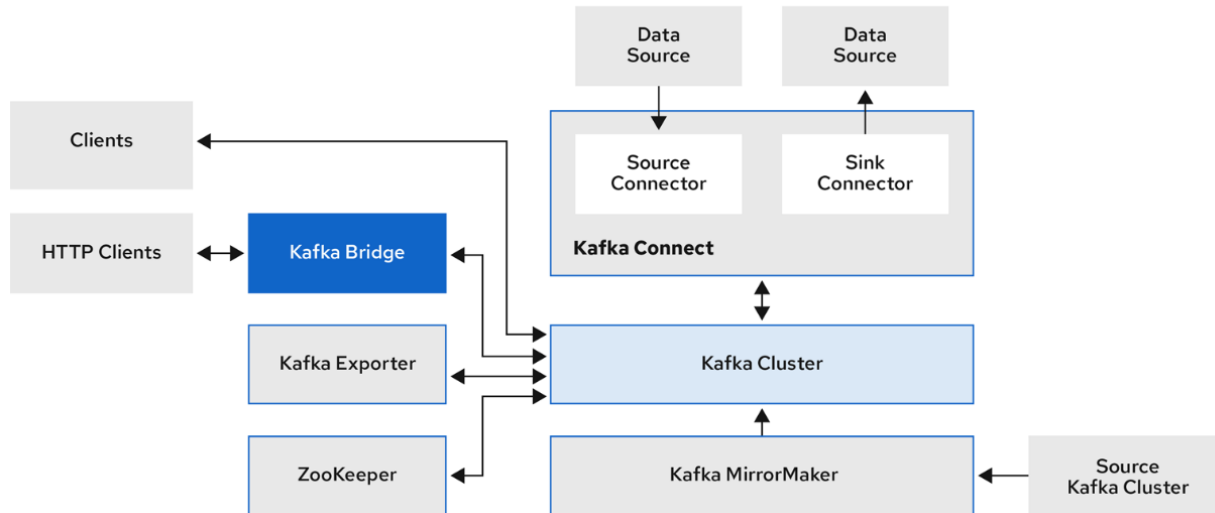


Technology backbone of Maersk

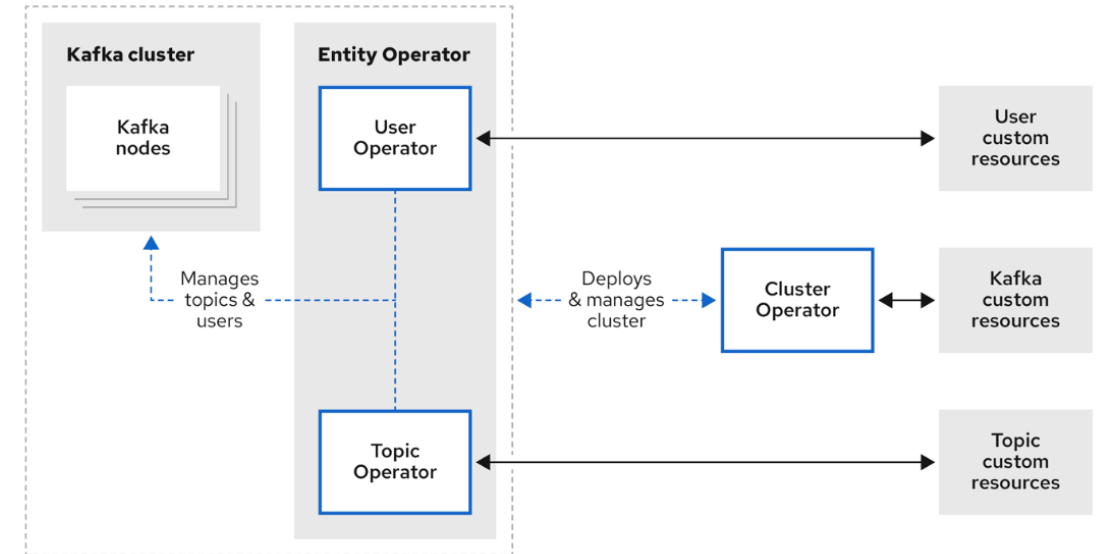


Running Kafka on Kubernetes

- Strimzi provides container images and operators for running Kafka on Kubernetes
- Strimzi operators are purpose-built with specialist operational knowledge to effectively manage Kafka on Kubernetes
- Apache Kafka components are provided for deployment to Kubernetes with the Strimzi distribution



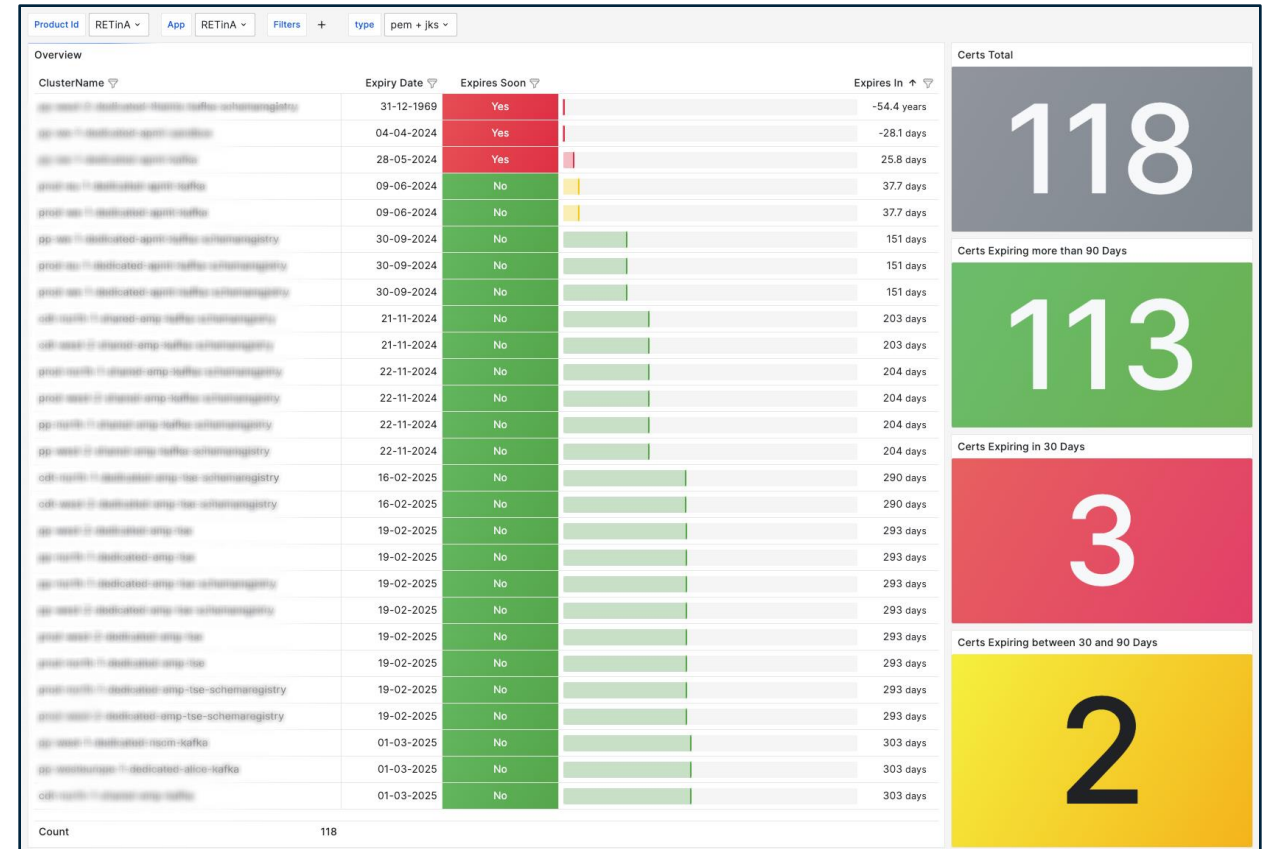
https://strimzi.io/docs/operators/latest/overview#overview-components_str



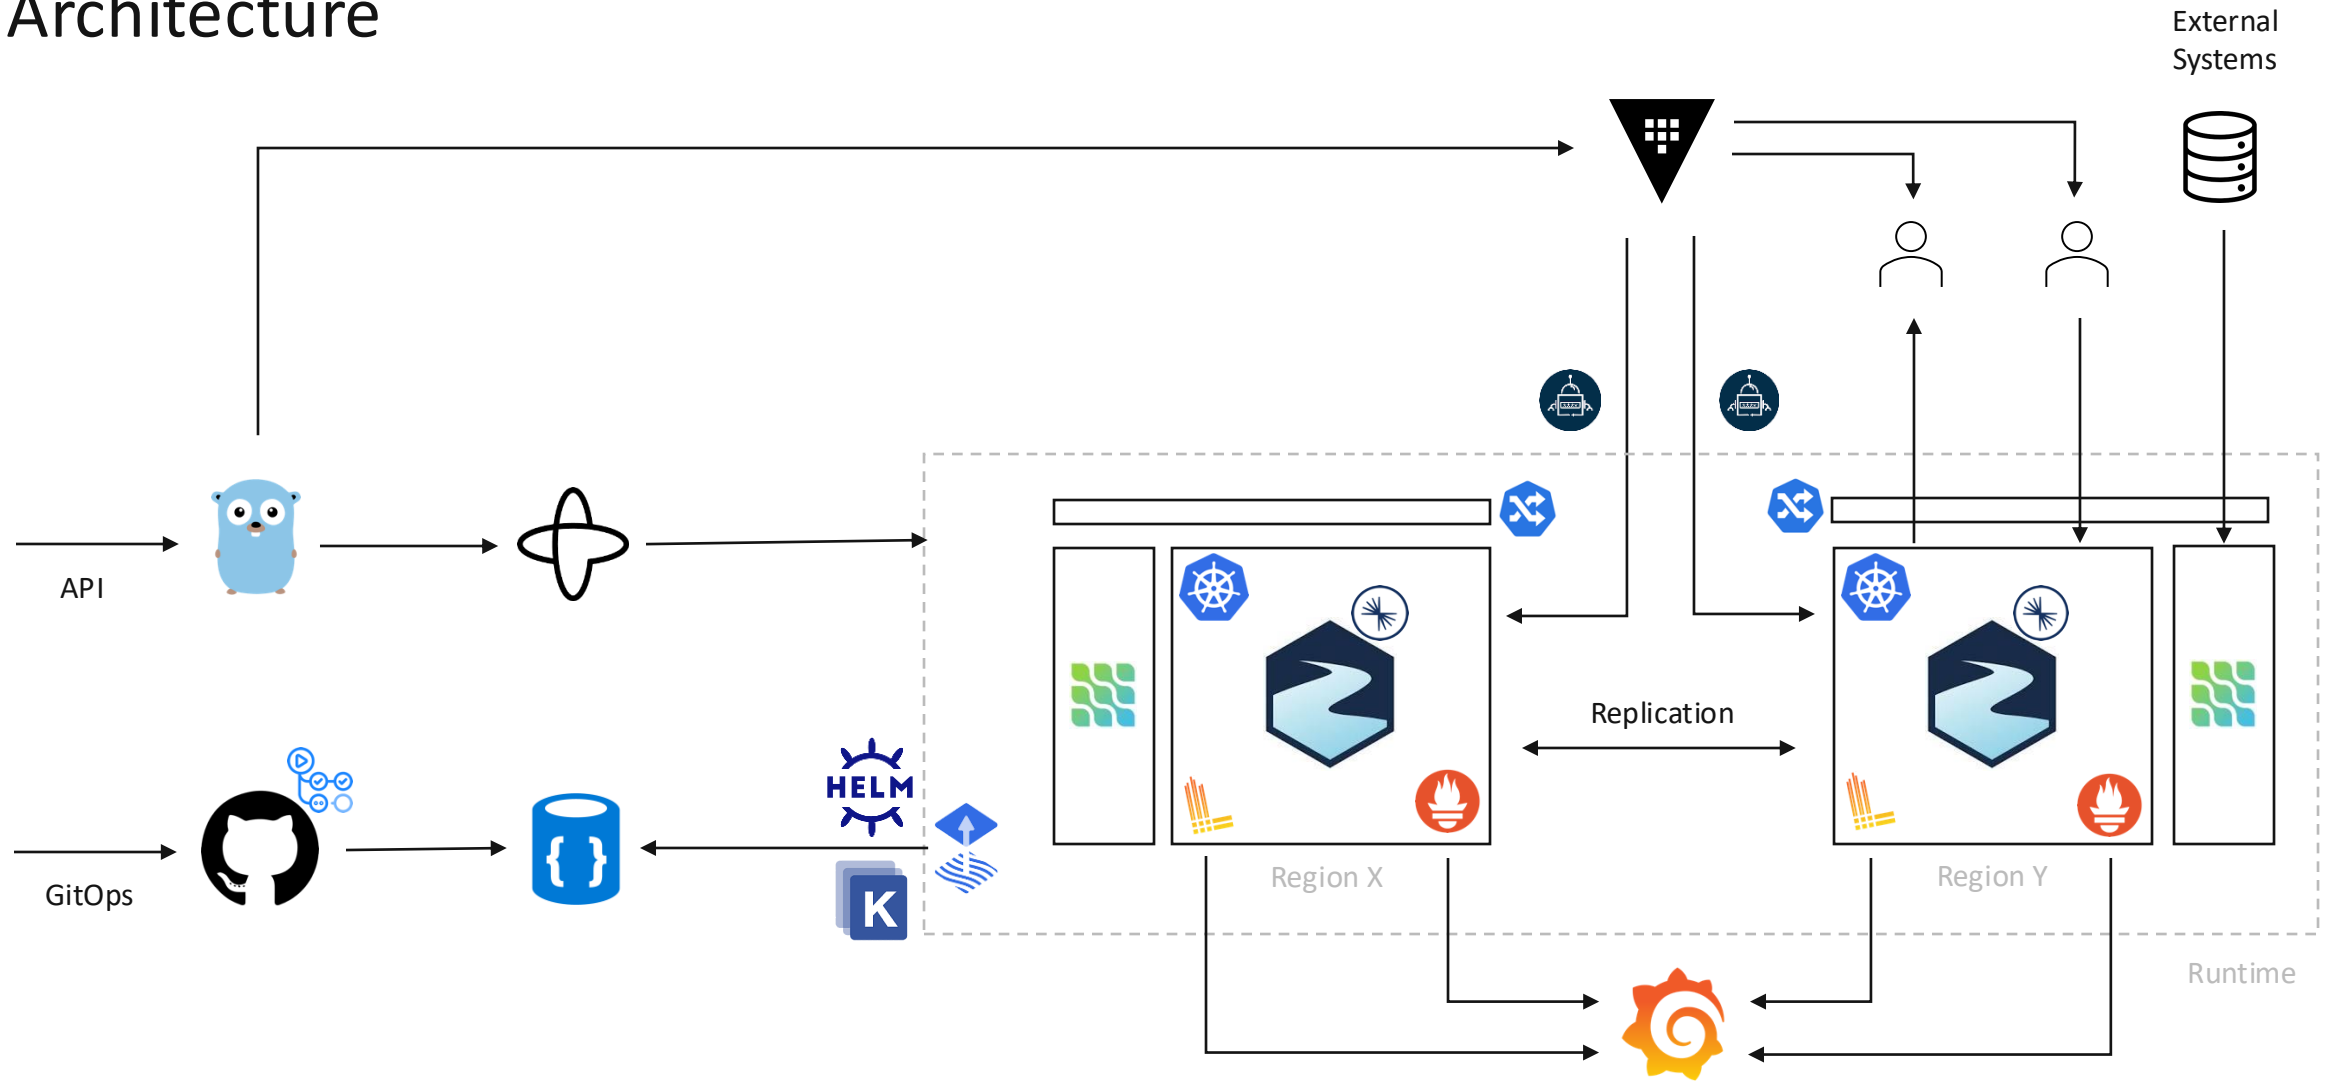
- Operators are a method of packaging, deploying, and managing Kubernetes applications
- An operator provide a way to extend the Kubernetes API and simplify the administration tasks associated with specific applications, i.e. Kafka Topic

Strimzi Contributions


- [#3761 Provide metrics to monitor certificates expiration](#)
- [#2779 CrdGenerator validate @JsonPropertyOrder](#)
- [#8732 Enhance KafkaBridge resource with consumer inactivity timeout and HTTP consumer/producer parts enablement](#)
- [#9537 Kafka Exporter Grafana dashboard too long URL error](#)
- [#7374 Improve the Kafka brokers Grafana dashboard](#)



Architecture



New API



OUR PURPOSE

Improving life for all by integrating the world



— The integration illustrated by five years of Automatic Identification System (AIS) transponder data from A.P. Moller - Maersk vessels registered in the company's scheduling system GSIS

● Gateway and hub terminals

Tenants

Number of teams using our Strimzi-based solution	+300
Number of Strimzi-based clusters actively running	81
Number of brokers in all our clusters	+450

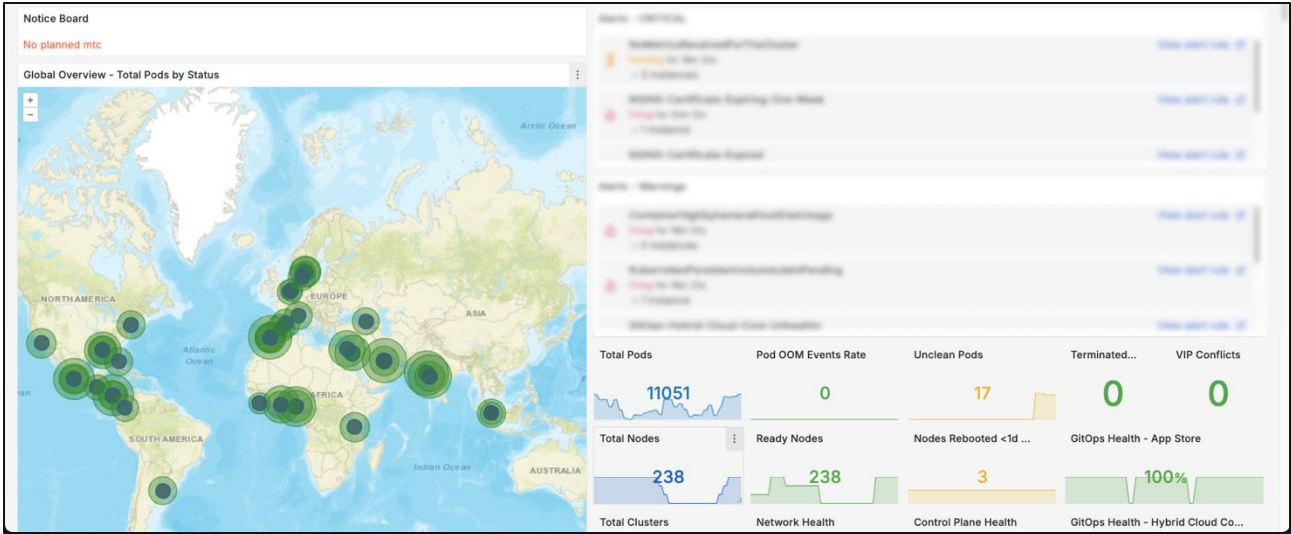
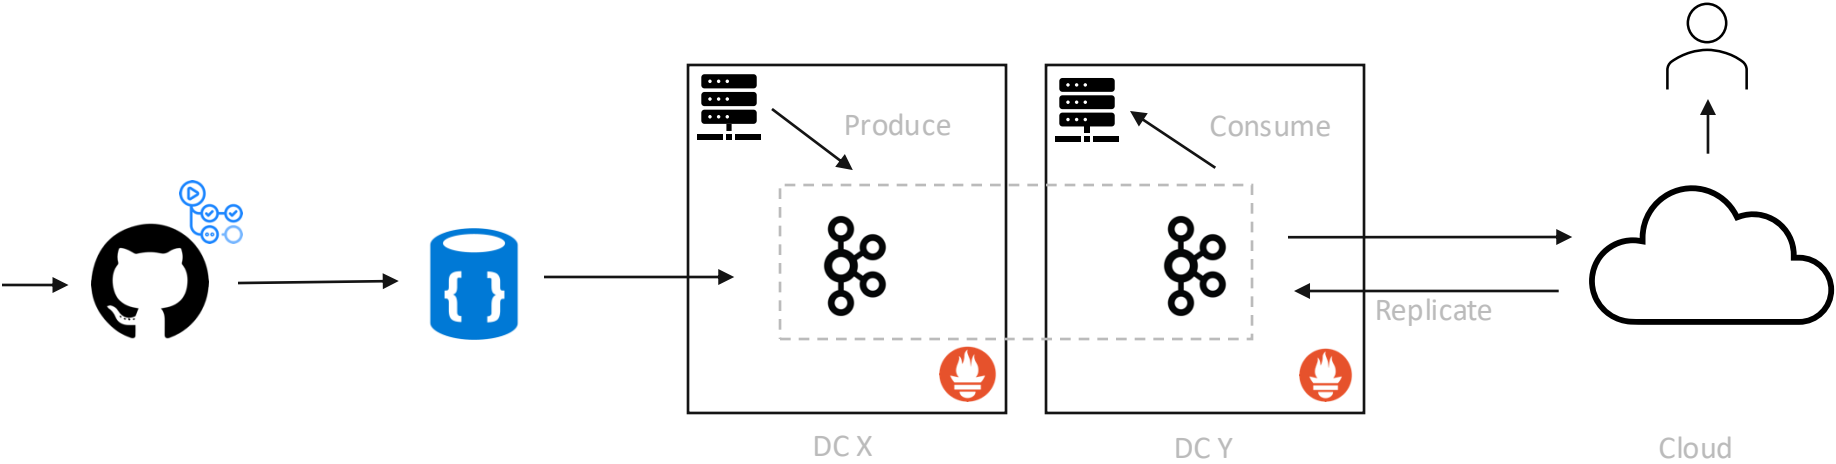
Topics

Total number of topics on all our clusters	+36K
Total number of partitions on all our topics	+410K
Number of AVRO and JSON schema versions in Schema Registry	+42K

Events

Total number of unique messages produced per day	+6Bil
Bytes in per second, avg	~140Mb
Bytes out per second, avg	~260Mb

On-Prem and Edge



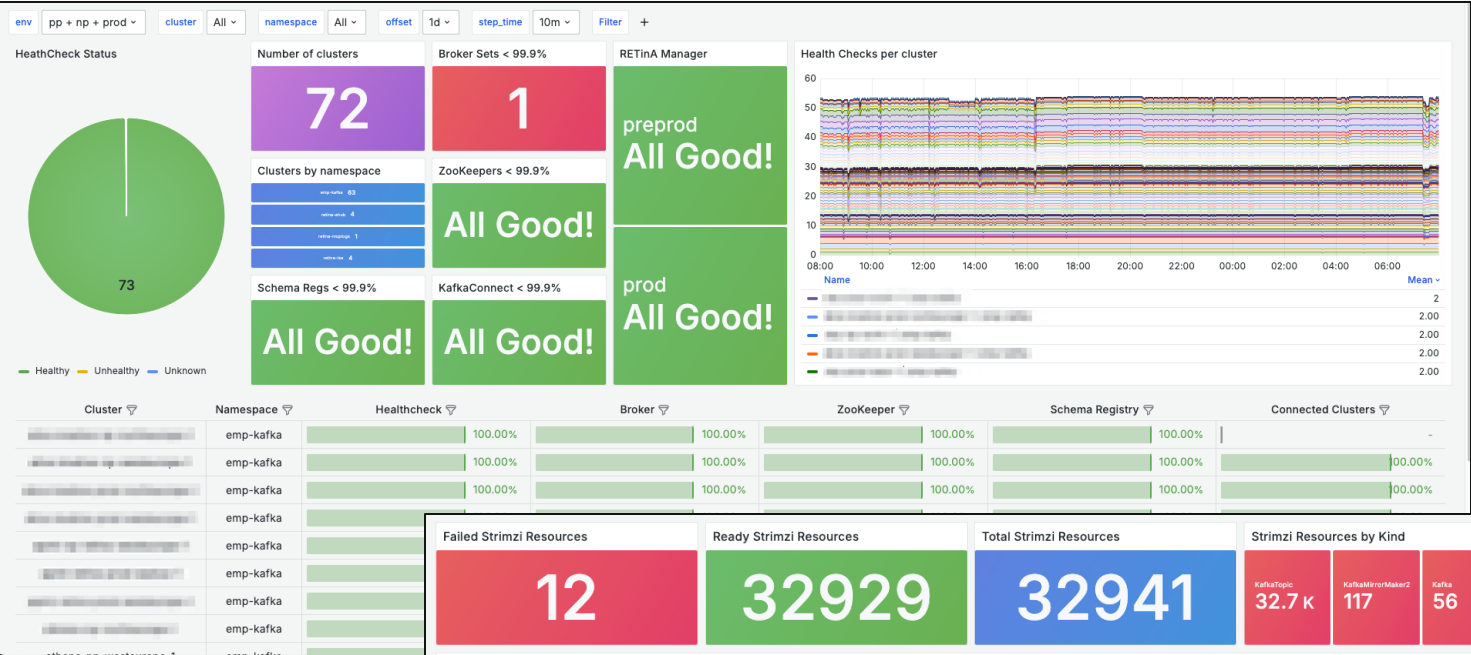
Edge

Number of topics in the terminals	+3.5K
Number of Strimzi-based clusters actively running in the terminals	37
Unique geographical locations	+17

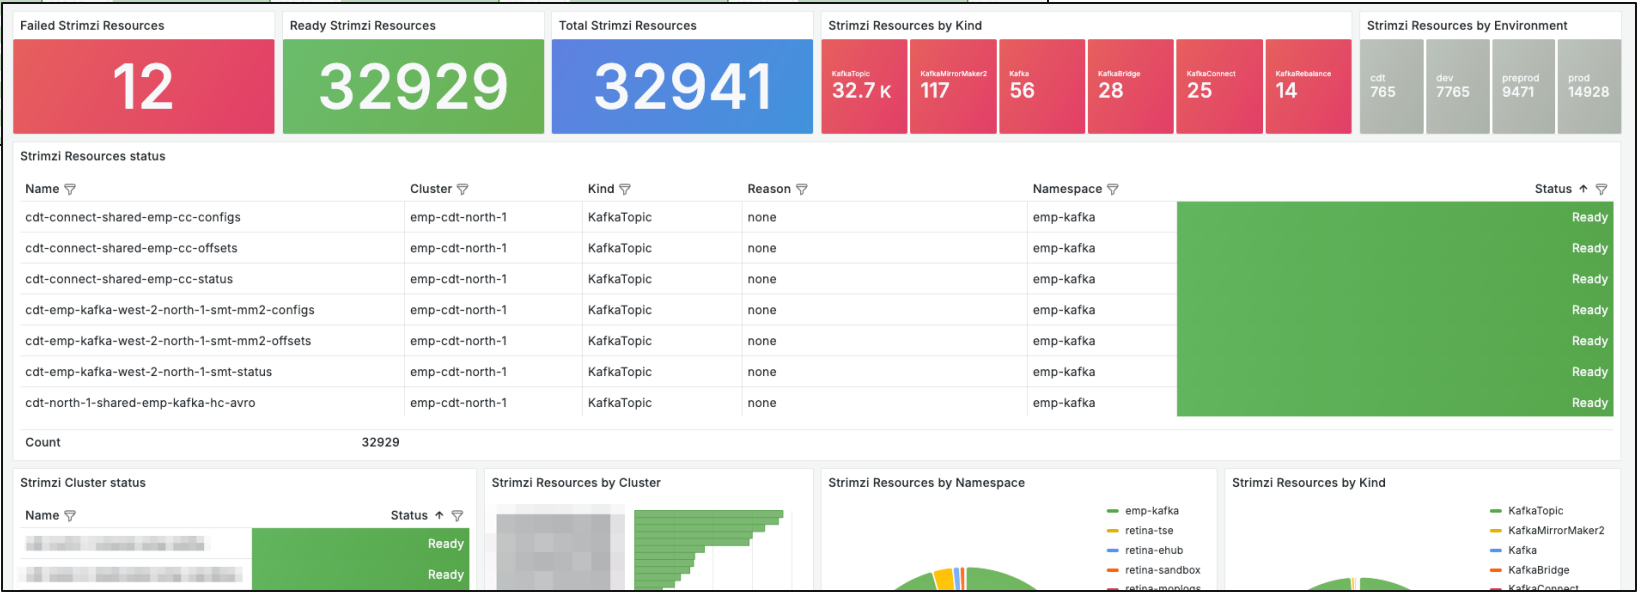
On-prem

Number of topics on-prem	+1.6K
Total number of unique messages produced per day	+100M
Total number of partitions on all our topics	+20K

Observability & Monitoring

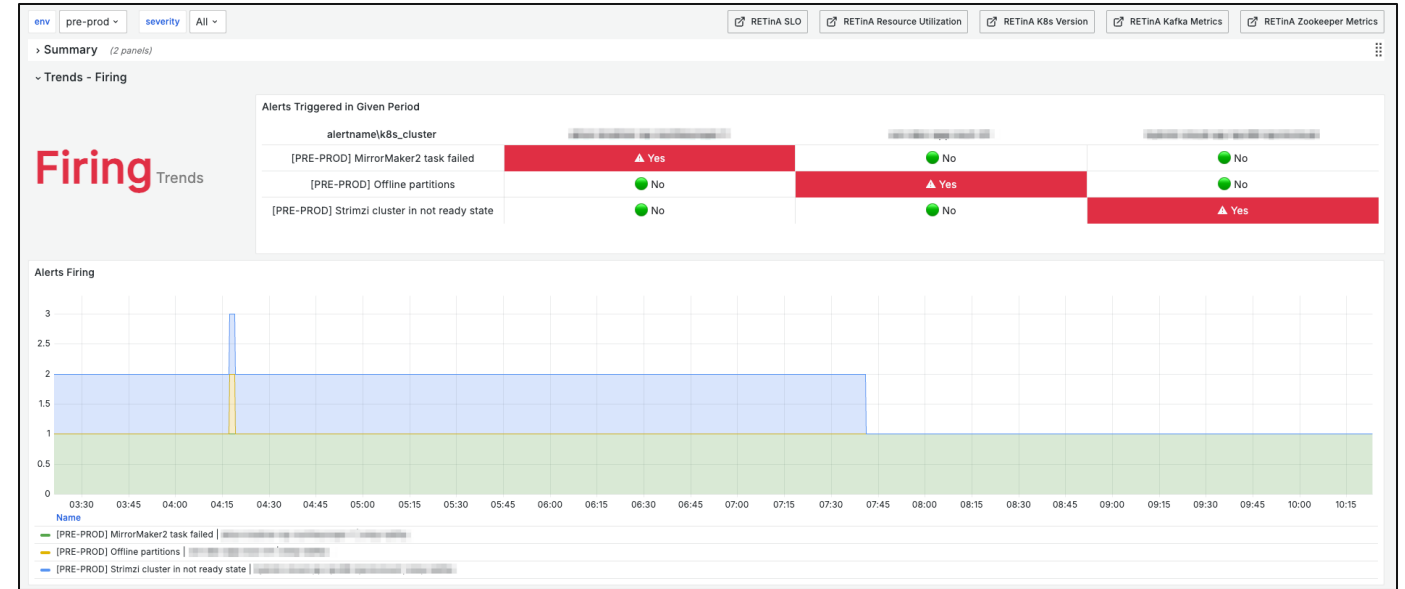


- Fully internal and open-source based observability platform
- Based on the LGTM and OpenTelemetry
- Healthchecks
- Real User Monitoring
- Synthetic Monitoring



Alerting

- Aggregated alerting by cluster and namespace
- Alerts and Dashboards as code
- Automated deployment and synchronization using GitHub Actions
- SLA/SLO 99.99%



[PRE-PROD] Disk will be full in 4h

One or more Disks will be full in 4h for emp-shared-pp-west-1 in namespace emp-kafka

Alert Id 355a047c-f9e6-4fd5-9838-9b4292f2293a

Description One or more Disks will be full in 4h for emp-shared-pp-west-1 in namespace emp-kafka.
Cluster: [REDACTED] Namespace: emp-kafka

Scope Name SC-PP-RETinA

Alert Status resolved

Fired At Wed Oct 16 10:25:09 UTC 2024

Alert Timestamp Wed Oct 16 10:34:11 UTC 2024

Labels

alert_id disk-full-in-4h

env pre-prod

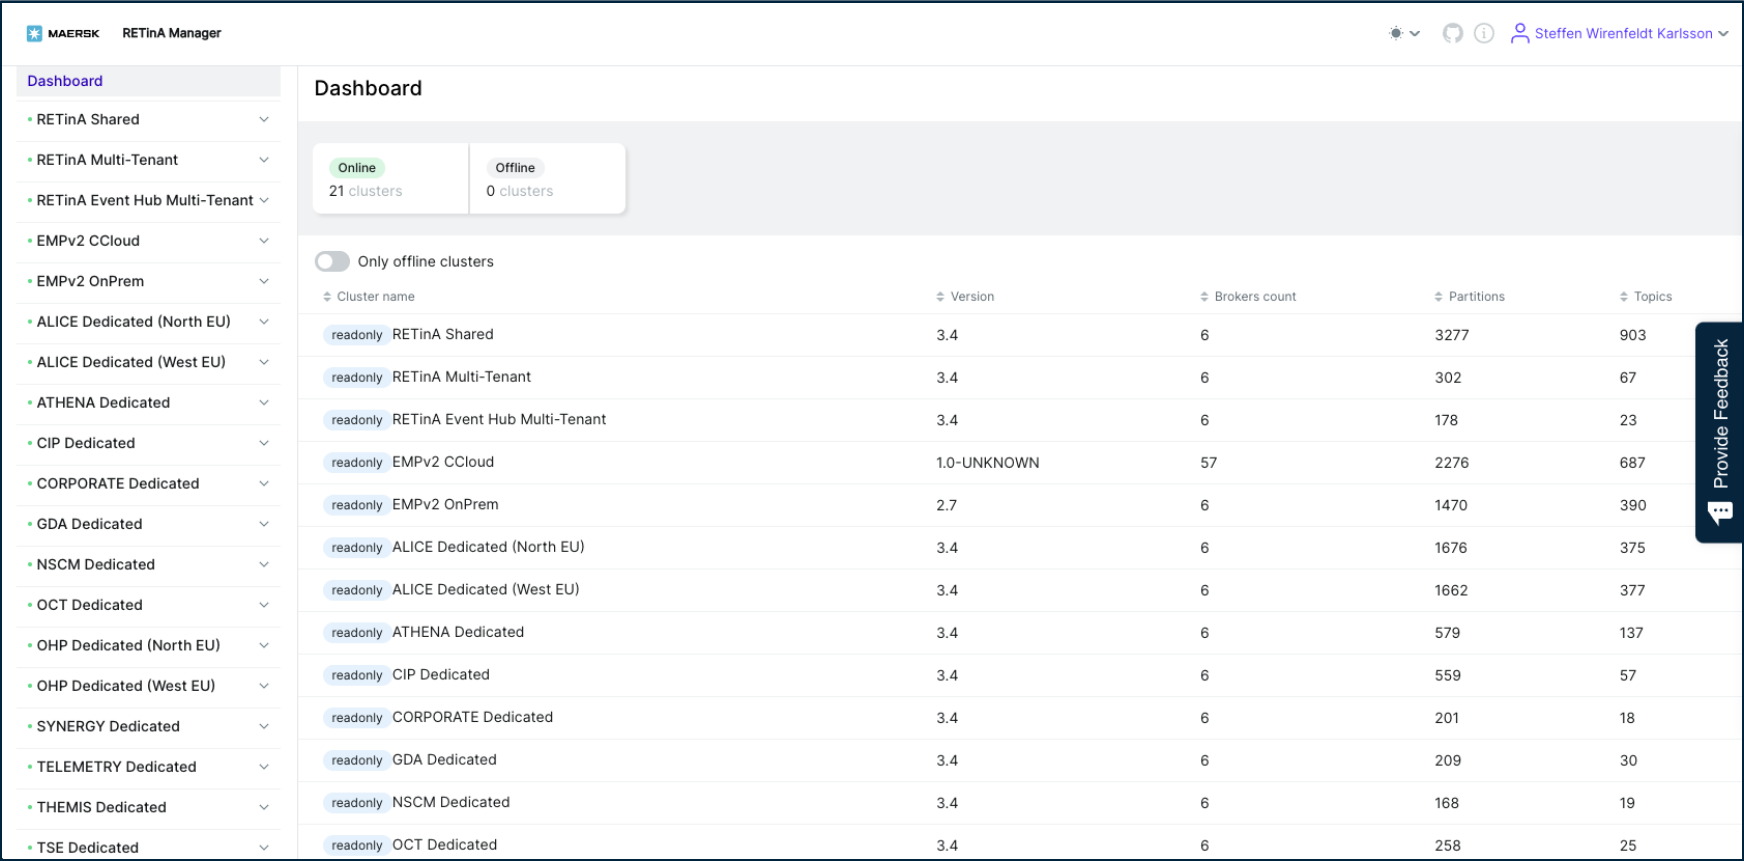
k8s_cluster [REDACTED]

namespace emp-kafka

priority P3 MI

[View alert](#) [Grafana](#)

RETinA Manager



Unique users
per day

+300



Requests
per day

+7.5K

Schema Compatibility UI

- Standalone and open-source Apache-2.0 license
- Fail fast and enable tenants to be more independent to decrease time-to-market
- Full self-service and transparency in schema compatibility and comparison
- **Schema Types:**
Avro, JSON and Protobuf
- **Compatibility levels:**
Backwards, Forwards, Full and None

The screenshot displays the Schema Compatibility UI interface. At the top, there are dropdown menus for 'Schema Type' (set to 'Avro') and 'Compatibility Level' (set to 'Backward'). Below these, two JSON schemas are shown side-by-side for comparison. The left schema is for 'Employee' with fields 'Name' (string) and 'Age' (int). The right schema is also for 'Employee' but includes an additional field 'Age2' (int). A red error message at the bottom left states: 'Validation failed: READER_FIELD_MISSING_DEFAULT_VALUE. The field Age2 at path /fields/2 in the new schema has no default value and is missing in the old schema, Age2'. A green 'VALIDATE' button is located at the bottom right.

```
1 {  
2   "type" : "record",  
3   "namespace" : "Tutorialspoint",  
4   "name" : "Employee",  
5   "fields" : [  
6     { "name" : "Name", "type" : "string" },  
7     { "name" : "Age", "type" : "int" }  
8   ]  
9 }
```

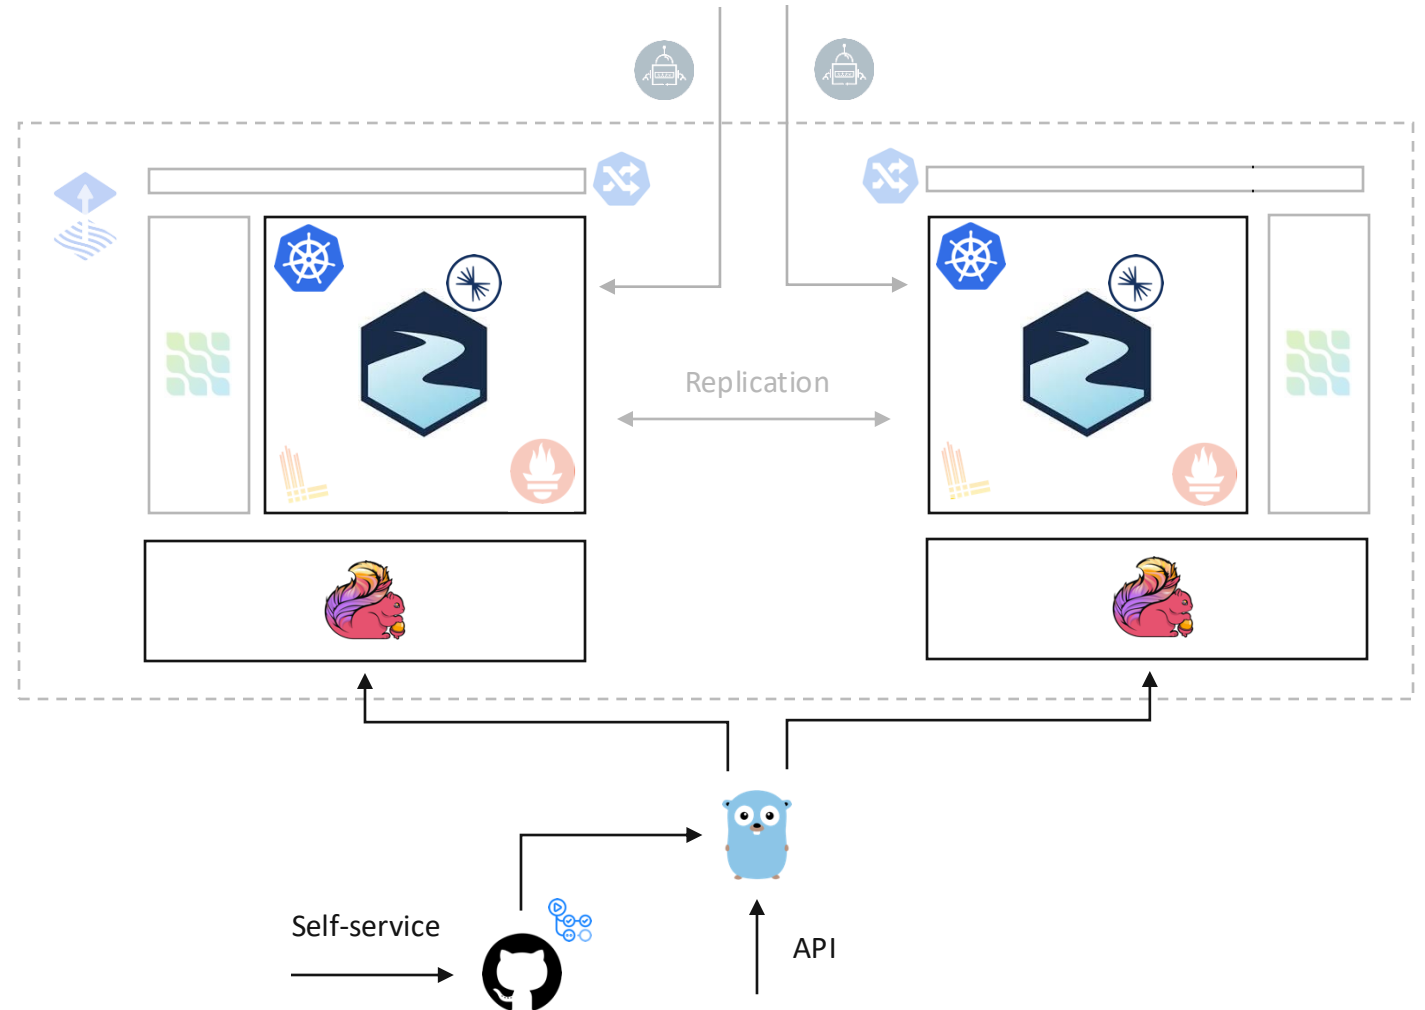
```
1 {  
2   "type" : "record",  
3   "namespace" : "Tutorialspoint",  
4   "name" : "Employee",  
5   "fields" : [  
6     { "name" : "Name", "type" : "string" },  
7     { "name" : "Age", "type" : "int" },  
8     { "name" : "Age2", "type" : "int" }  
9   ]  
10 }
```

Validation failed:
READER_FIELD_MISSING_DEFAULT_VALUE
The field Age2 at path /fields/2 in the new schema has no default value and is missing in the old schema, Age2

VALIDATE

Stream Processing

- Automated self-serviceable stream processing on-top of Strimzi with Flink and GitHub Actions
- API-first architecture for better system-to-system integrations
- Predefined template jobs for better overall performance on the cluster
- Deployed using open-source community operator for better stability and configuration of Flink
- Fully monitored and observed using the Grafana and Prometheus stack



Reflections

- Migrations are hard in kafka world
- More clusters = more problems
- Supporting multiple onboarding scenarios ain't easy
- Good observability ease support pain
- In our case migrating from confluent cloud cause significant savings even with more adoption of current platform
- Operational stability is a key
- Supporting multiple deployment environments (edge, onprem, cloud) requires wide knowledge spread across team

